




Real-Time Operating System Basics

CSE4354/5354





Scope

- **This is a very simplified survey of real-time operation system (RTOS) operation and computer science concepts, condensed to fit into two lecture periods**
 - **This document is intended to show a minimum set of concepts required to complete the class projects**
 - **It is recommended that you take an advanced computer sciences OS course if designing or using a RTOS**
- 

References


- **G. L. Peterson. *Myths about the mutual exclusion problem*. Information Processing Letters, 12(3):115--116, 1981**
- **Eisenberg, M.A., and McGuire, M.R. *Further comments on Dijkstra's concurrent programming control problem*. Comm. ACM 15, 11 (Nov. 1972), 999**
- **Knuth, D.E.: *The art of computer programming*. Fundamental algorithms. Addison-Wesley, 3rd edition (1997)**

References

- **Free Dictionary of Online Computing**
 - <http://wombat.doc.ic.ac.uk/foldoc/>
- **Scheduling (MUF and Quick Review of RM, EDF, and MLF)**
 - <http://www.ee.umd.edu/serts/bib/book/article/rtp92.shtml>
- **Priority inversion on Mars Pathfinder**
 - http://research.microsoft.com/~mbj/Mars_Pathfinder/




RTOS Topics

- **What is a RTOS?**
 - **Tasks, Processes/Threads, and Kernel**
 - **Multitasking**
 - **Scheduling**
 - **Critical Sections**
 - **Inter-process Communication**
- 




What is a RTOS?

- **A real-time operating system is one in which a set of computing tasks are considered correct only if the tasks are completed in a timely manner**
 - **A real-time operating system executes these tasks correctly in a predictable (deterministic) manner**
- 



Common OS Terms


- **Processes** - Complicated tasks require “heavy weight” code that requires a lot of state information and does not share a memory space with other processes
 - **Thread (of execution)** - Simple tasks have “light weight” code that shares the memory space within the process
 - **Kernel** - A part of the operating system responsible for managing system resources, scheduling and dispatching processes, and inter-process communications
- 

Task Properties

- **Periodicity (Synchronicity)**
 - Task can be timer-driven (periodic/synchronous)
 - Task can be event-driven (aperiodic/asynchronous)
 - Task can be a one-shot event (special aperiodic case)
- **Temporal Attributes**
 - Task duration (how long to complete a task)
 - Start deadline (must start by some time)
 - Stop deadline (must complete by some time)
- **Deadline Type (hard or soft deadline)**
- **Relative Priority**



Tasks

- **A task has at least three states**
 - **Running, blocked, and ready**
 - **Only one task can run at a time on the M4F controller (no parallel processing)**
 - **Scheduler allows concurrent tasks using multi-tasking techniques (pseudo-parallelism)**
 - **A task control block (TCB) keeps track of the state of a thread**
- 

Task Execution

- **Task can be started by**
 - **Interrupt**
 - **Scheduler decisions**
- **Execution of interrupt task**
 - **Can run to completion**
 - **Can be preempted by a higher-priority interrupt**
- **Execution of dispatched tasks**
 - **Can run to completion (batch-like operation)**
 - **Can autonomously relinquish control (cooperative multi-tasking) before completion**
 - **Can be preempted by the scheduler (preemptive multi-tasking)**

Kernel Scheduler

- **Must decide which task to run**
- **Handling periodic tasks is straightforward**
- **There is no apriori knowledge of when asynchronous events will occur, so the kernel must rely on statistics to determine scheduling (min time between tasks, average time between tasks, ...)**
- **May also consider the timing and type of each task deadline to prioritize scheduling (hard deadlines have catastrophic consequences if missed)**

Kernel Dispatcher

- **Switches the context from a task to a new task selected by the scheduler**
- **Context switch steps**
 - **If preemptive multitasking, stops any currently running task if applicable and saves the context**
 - **Switches the context to the scheduler selected process (if restarting, then restores the context)**
 - **On the M4F, the context consists of the R0-12, SP, LR, PC, xPSR, S0-15, FPSCR, and other registers**

Multi-tasking

- **Cooperative**
 - Kernel lets a task execute until completion (may temporarily interrupt it for interrupt processing)
 - If task is short, task runs to completion
 - If task is long, the task cooperates by relinquishing control back to the kernel so that other tasks can run
 - “Long” and “short” determine cooperation
 - Problem is a task “crashes” and does not relinquish control to other tasks
- **Preemptive**
 - Kernel assigns a quantum (time slice) to the task
 - If the task does not complete within the slice, the kernel preempts the task and lets another run

Scheduling 101


- **Schedule should be feasible**
- **Need to avoid task starvation**
- **Need to watch for deadlock (or prevent it)**
- **Sometimes, a watchdog timer is used to detect “lockups” in the system**
- **Latency of dispatcher and scheduler to should be minimized**
- **Should optimize context switching time**
 - **Too fast causes thrashing (too much managing and no working)**
 - **Too slow causes sluggishness**

Non-Priority Scheduling


- **FIFO**
 - First task into a queue executes until completion
- **Round robin**
 - Preemption method where N tasks are each dispatched in order with equal quantum
- **Rate Monotonic (RM)**
 - Shortest pending task is executed first

Priority Scheduling

- **Earliest Deadline First (EDF)**
 - The most critical deadline (function of whether the deadline is hard or soft and the time to expiration) determines who goes first
- **Minimum Laxity First (MLF)**
 - Like EDF, but prioritization based on laxity
 - $\text{Laxity} = \text{time to deadline} - \text{execution time}$
- **Maximum Urgency First (MUF)**
 - An RM schedule with a constraint that some tasks are temporarily starved if CPU cycles are not available



Critical Sections

- **When multiple tasks access the same hardware resource or shared memory conflicts can occur**
 - **When a task needs exclusive access to a resource, it enters a critical section of code, where if another task interferes (collides), the result of the operation could be corrupted**
 - **Mutexes, semaphores, and inter-process signaling can be used to prevent collisions by blocking entry into a critical section in use**
- 

Critical Sections

- Sub-optimal solution could be to disable interrupts for a few clocks
- Could only be used on short critical sections
- Can have dire impacts to system performance – bad practice
- Code snippet:

```
disable_interrupts (GLOBAL) ;  
    critical section  
enable_interrupts (GLOBAL) ;
```
- We will use alternative techniques

Inter-process (thread) Signaling

- Can be used to coordinate access to shared resources
- Can synchronize multiple tasks so that a task only runs when other data from other tasks are ready
 - Called the producer-consumer scenario
 - Can coordinate reading and writing data from keyboard, uart, ... through a ring buffer in shared memory

Mutual Exclusion

- **Mutual exclusion object (mutex) coordinates access to a shared resource as follows:**
- **As an atomic (indivisible) operation, the requesting task**
 - **Checks the status of a mutex**
 - **If unlocked, it**
 - **Locks the mutex**
 - **Enters the critical section**
 - **Unlocks the mutex**
- **If the mutex was locked, access is blocked and the task waits**
- **It could use busy waiting (spinlock), waiting in a queue, or a combination**

Priority Inversion

- **A case where a lower priority task runs instead of a higher priority task**
- **Example with mutex operation:**
 - Low, medium and high priority tasks (L, M, and H)
 - L is running and locks a mutex that H needs
 - H starts running and is blocked
 - L would normally run and unlock the mutex, but M preempts L, leaving H blocked from running
- **This example could be solved through priority inheritance (temporary elevation of L process priority to H process priority when H process is blocked waiting on the locked mutex)**

Errant Mutual Exclusion Solution

- Not atomic (another task/ISR could execute code between the while (lock) {} and lock = TRUE code)
- Conceptual code snippets:

```
// init
bool lock = FALSE; // global

...

// task code
while (lock)
    {"busy wait (spinlock)"}
lock = TRUE;
    critical section
lock = FALSE;
```

Errant Interrupt Mutual Exclusion Solution

- This solution is errant, since if locked, the system will lockup
- Conceptual code snippets:

```
// init
bool lock = FALSE; // global
...
// process code
disable_interrupts(GLOBAL);
while (lock) {"wait"}
lock = TRUE;
enable_interrupts(GLOBAL);
    critical section
lock = FALSE;
```

Interrupt Masking Mutex Solution

- Interrupts make atomic on single core processors, but could degrade performance or cause interrupt loss - bad practice
- Conceptual code snippets:

```
// init
bool lock = FALSE; // global

...
// task code
bool ok = FALSE;
while (!ok)
{
    disable_interrupts(GLOBAL or TASK_SWITCH_ISR);
    if (!lock) {ok = lock = TRUE;}
    enable_interrupts(GLOBAL or TASK_SWITCH_ISR);
}

    critical section
lock = FALSE;
```


Software Mutex Solutions

- **Dekker**
 - First software-only solution with 2 tasks
- **Peterson**
 - Simpler than Dekker's solution with 2 tasks
- **Eisenberg and McGuire**
 - Simple N task solution with rotating priorities

Peterson's Mutex Solution

- Two task solution, simpler than Dekker's algorithm
- Code snippets:

```
bool busy[2] = {FALSE, FALSE}; // global
int turn = 0; // global

...


// task i code; i = 0..1
busy[i] = TRUE;
turn = 1-i;
while (busy[1-i] && turn != i) { };
    critical section
busy[i] = FALSE;
```

Peterson's Mutex Solution

- Suppose task 0 is waiting to enter the critical section (executing the while loop)
- 2 conditions allow entry into the critical section:
 - Task 1 not busy (trying to use the resource or using the resource)
 - OR it is task 0's turn
- The critical section in task 0 can run if the last code task 1 executed was
 - Before `busy[i] = TRUE` or after `busy[i] = FALSE` (other task is not busy)
 - After `turn = 1 - i` (turn given away by process 1) (task 0's turn)



Hardware Mutex Solution

- **On some processors, there is hardware support for atomic operations on memory**
 - **For instance, on the single core M4F, there are variants of the load and store commands, LDREX and STREX, that allow locking and exclusive access to memory**
- 

Semaphores

- **A synchronization method that shares a collection of resources across tasks**
- **Binary semaphore are 0 or 1 valued**
- **Counting semaphores have values ≥ 0**
- **Functions (three different nomenclatures) `wait()`, `down()`, or `P()`**
 - **Puts task to sleep and adds task to a queue if count = 0**
- **Decrements count and allow the task to proceed if count > 0**
 - `post()`, `signal()`, `up()`, or `V()`**
 - **Increments count**
 - **If count = 1 after incrementing, then a task may be waiting to run in the queue**
 - **If task waiting in queue, wake up the task**
- **Generally, a task either waits or posts, not both**

Semaphore Implementation

- **Conceptual view of a semaphore structure**
- **Code snippet**

```
typedef struct _semaphore
{
    int count;
    int queue_count;
    int process_queue[MAX_SIZE];
} semaphore;
```

Semaphore Functions

- **Conceptual Atomic Functions**

- **wait(semaphore& s)**

```
while (s.count == 0)
{ s.task_queue[queue_count++] =
  current_task;

  dispatch other task(s)

  task(s) post the semaphore and queues up waiting
  task(s), prog flow eventually returns }

s.count--;
```

- **post(semaphore& s)**

```
s.count++;

if (s.count == 1)
{ if (queue_count > 0)
  select task to run, dec s.count,
  set state to READY state,
  return to caller }
```

Producer-Consumer Example

- Consider a solution with three synchronization primitives
 - One mutex (key) used to control access to inventory
 - One counting semaphore (needed) used to control producing
 - One counting semaphore (available) used to control consuming
- Maximum number of inventory items is N

Semaphore Example

- **Code snippet**

```
mutex key = false;  
semaphore needed, available;  
needed.count = N;  
available.count = 0;
```

Semaphore Example

- **Code snippet**

```
// producer task
wait(needed) ;
lock(key) ;
    add widget
unlock(key) ;
post(available) ;
...
// consumer task
wait(available) ;
lock(key) ;
    remove widget
unlock(key) ;
post(needed) ;
```

Common Bottlenecks for RTOS

- **Serving**
 - Polling overhead
(use interrupts or semaphores when possible to mitigate)
 - Latency of interrupts may affect performance
- **Multi-tasking**
 - Latency of context switching
(save as little as absolutely possible)
 - Impact of blocking
(use sleep and wake instead of busy waiting when possible)
- **Computational**
 - Complex mathematical functions
(use lookup tables when possible)